

Securing your daemons using systemd

Zbigniew Jędrzejewski-Szmek



zbyszek@in.waw.pl



Brno, 2020.1.26

Before we begin...

Why use systemd for this at all?

- ▶ centralization
- ▶ abstraction of hardware architecture / kernel version
- ▶ unprivileged operation

Before we begin...

Why use systemd for this at all?

- ▶ centralization
- ▶ abstraction of hardware architecture / kernel version
- ▶ unprivileged operation

Numbers:

```
$ dnf repoquery --releasever=32 -l --whatprovides \  
    '/usr/lib/systemd/system/*' \  
    rg -i '^/usr/lib/systemd/system/[a-z0-9_@.\-]+$' | \  
    sort -u | wc -l  
1740!!!
```

Before we begin...

Unit files

Before we begin...

Unit files

```
# /etc/systemd/system/mydaemon.service
```

```
[Service]
```

```
ExecStart=/usr/local/bin/mydaemon
```

```
$ sudo systemctl start mydaemon.service
```

Before we begin...

Unit files

```
# /etc/systemd/system/mydaemon.service
```

```
[Service]
```

```
ExecStart=/usr/local/bin/mydaemon
```

```
$ sudo systemctl start mydaemon.service
```

```
$ sudo /usr/local/bin/mydameon
```

Before we begin...

Unit files

```
# /etc/systemd/system/mydaemon.service
```

```
[Service]
```

```
ExecStart=/usr/local/bin/mydaemon
```

```
$ sudo systemctl start mydaemon.service
```

```
$ sudo /usr/local/bin/mydameon
```

```
$ sudo systemd-run /usr/local/bin/mydameon
```

```
$ sudo systemd-run -t /usr/local/bin/mydameon
```

Basics

User=

Basics

User=

```
$ systemd-run whoami  
root
```

```
$ systemd-run -p User=zbyszek whoami  
zbyszek
```

```
$ systemd-run --uid=zbyszek whoami  
zbyszek
```

Limiting access to the file system

- ▶ `ProtectHome=yes|read-only`
- ▶ `ProtectSystem=yes|full|strict`

Limiting access to the file system

- ▶ `ProtectHome=yes|read-only`
- ▶ `ProtectSystem=yes|full|strict`
- ▶ `InaccessiblePaths=`
- ▶ `ReadOnlyPaths=`
- ▶ `ReadWritePaths=`

Limiting access to the file system

- ▶ `ProtectHome=yes|read-only`
- ▶ `ProtectSystem=yes|full|strict`
- ▶ `InaccessiblePaths=`
- ▶ `ReadOnlyPaths=`
- ▶ `ReadWritePaths=`
- ▶ `BindPaths=`
- ▶ `ReadOnlyBindPaths=`

Limiting access to the file system

- ▶ `PrivateTmp=yes`

Limiting access to the file system
a better way

Limiting access to the file system a better way

- ▶ `RuntimeDirectory=foo` `/run/foo/`
- ▶ `StateDirectory=foo` `/var/lib/foo/`
- ▶ `CacheDirectory=foo` `/var/cache/foo/`
- ▶ `LogsDirectory=foo` `/var/log/foo/`
- ▶ `ConfigurationDirectory=foo` `/etc/foo/`

```
$ sudo systemd-run -t -p User=zbyszek \  
-p RuntimeDirectory=foo \  
ls -ld /run/foo
```

- ▶ automatic *creation* and *ownership*
- ▶ automatic *removal*

User creation on demand?

User creation on demand?

- ▶ `DynamicUser=yes`

User creation on demand?

► DynamicUser=yes

```
$ systemd-run -p DynamicUser=1 -t whoami
```

User creation on demand?

► DynamicUser=yes

```
$ systemd-run -p DynamicUser=1 -t whoami
```

```
$ echo -e 'asdf\nasdf' | \
```

User creation on demand?

► DynamicUser=yes

```
$ systemd-run -p DynamicUser=1 -t whoami
```

```
$ echo -e 'asdf\nasdf' | \
```

```
systemd-run --pipe -p DynamicUser=1 \  
    bash -c 'grep .; whoami' | \
```

User creation on demand?

► DynamicUser=yes

```
$ systemd-run -p DynamicUser=1 -t whoami
```

```
$ echo -e 'asdf\nasdf' | \
```

```
systemd-run --pipe -p DynamicUser=1 \  
    bash -c 'grep .; whoami' | \
```

```
systemd-run --pipe -p DynamicUser=1 \  
    bash -c 'grep .; whoami' | \
```

User creation on demand?

► DynamicUser=yes

```
$ systemd-run -p DynamicUser=1 -t whoami
```

```
$ echo -e 'asdf\nasdf' | \
```

```
systemd-run --pipe -p DynamicUser=1 \  
  bash -c 'grep .; whoami' | \  
systemd-run --pipe -p DynamicUser=1 \  
  bash -c 'grep .; whoami' | \  
systemd-run --pipe -p DynamicUser=1 \  
  bash -c 'grep .; whoami'
```

What about the network?

What about the network?

- ▶ `PrivateNetwork=yes`

What about the network?

- ▶ `PrivateNetwork=yes`

“`PrivateNetwork=yes` is the recommended way to run network services”

Socket Activation

A daemon does not open a socket itself, it receives a socket from the manager

Socket Activation

A daemon does not open a socket itself, it receives a socket from the manager

Two types of socket activation:

Accept=yes

→ a single instance of the service is started for each connection
→ “wait” under inetd/xinetd

Accept=no

→ a single instance of the service is started for each connection
→ “nowait” under inetd/xinetd

Low level stuff

- ▶ `MemoryDenyWriteExecute=yes`
- ▶ `PrivateDevices=yes`
- ▶ `NoNewPrivileges=yes`
- ▶ `RestrictAddressFamilies=AF_UNIX|AF_INET|AF_INET6
|AF_CAN|AF_APPLETALK|...`
- ▶ `RestrictSUIDSGID=yes`
- ▶ `ProtectKernelTunables=yes`
- ▶ `ProtectHostname=yes`
- ▶ `ProtectKernelLogs=yes`
- ▶ `SystemCallArchitectures=native|x86_64|i386|...`
- ▶ `LockPersonality=yes`

Capability limits

- ▶ `CapabilityBoundingSet=`
- ▶ `Capability=`
- ▶ `DropCapability=`
- ▶ `AmbientCapabilities=`

System call filtering

“seccomp mode 2”

System call filtering

“seccomp mode 2”

- ▶ `SyscallFilter=...`
implemented using `libseccomp`
- ▶ `syscall1 | syscall2 | @group`
- ▶ `@basic-io`
- ▶ `@obsolete`

System call filtering

“seccomp mode 2”

- ▶ SyscallFilter=...
implemented using libseccomp
- ▶ syscall1 | syscall2 | @group
- ▶ @basic-io
- ▶ @obsolete

```
$ systemd-analyze syscall-filter @obsolete
```

Per-service network firewall

- ▶ `IPAddressAllow=10.20.30.0/24 1.2.3.4`
- ▶ `IPAddressDeny=*`
- ▶ `IP{Ingress,Egress}FilterPath=`

systemd-analyze security

```
$ systemd-analyze security systemd-resolved.service
```

The End

`https://github.com/systemd/systemd`

docs: `https://systemd.io/`

`https://www.freedesktop.org/wiki/Software/systemd/`

this: `https://github.com/keszybz/systemd-security-talk`

`https://github.com/keszybz/jesien-systemd-security/`

`blob/master/systemd-security.pdf`